

# Distributional Semantics


COMP90042

Natural Language Processing

Lecture 10



THE UNIVERSITY OF  
MELBOURNE



# Lexical Databases - Problems

- Manually constructed
  - ▶ Expensive
  - ▶ Human annotation can be biased and noisy
- Language is dynamic
  - ▶ New words: slang, terminology, etc.
  - ▶ New senses
- The Internet provides us with massive amounts of text. Can we use that to obtain word meanings?

# Distributional Hypothesis

- “You shall know a word by the company it keeps” (Firth, 1957)
- Document co-occurrence often indicative of topic (*document as context*)
  - ▶ E.g. *voting* and *politics*
- Local context reflects a word’s semantic class (*word window as context*)
  - ▶ E.g. *eat a pizza*, *eat a burger*

# Guessing Meaning from Context

- Learn unknown word from its usage
- E.g., tezgüino*
  - (14.1) A bottle of \_\_\_\_\_ is on the table.
  - (14.2) Everybody likes \_\_\_\_\_.
  - (14.3) Don't have \_\_\_\_\_ before you drive.
  - (14.4) We make \_\_\_\_\_ out of corn.
- Look at other words in same (or similar) contexts

	(14.1)	(14.2)	(14.3)	(14.4)	...
<i>tezgüino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	0	

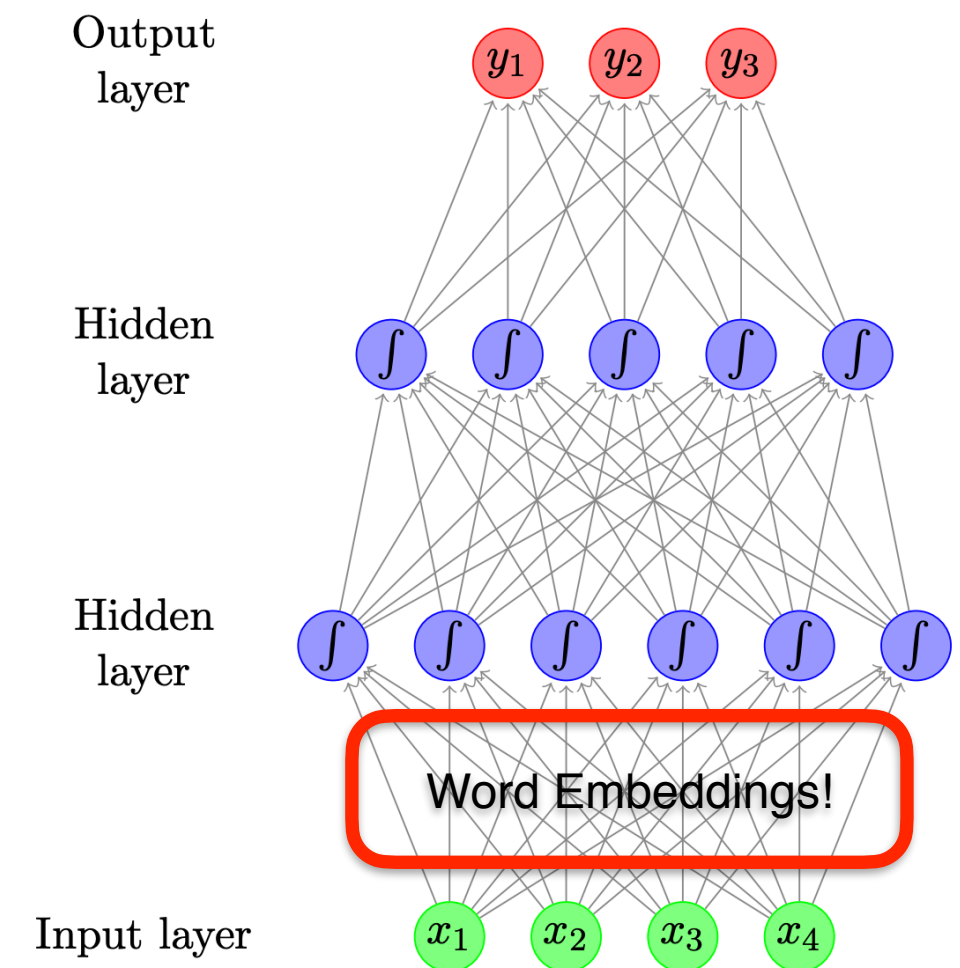
# Word Vectors

	(14.1)	(14.2)	(14.3)	(14.4)	...
<i>tezgiino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	0	

- Each row can be thought of a **word vector**
- It describes the distributional properties
- Capture all sorts of semantic relationships (synonymy, analogy, etc)

# Word Embeddings?

- We've seen word vectors before: word embeddings!
- Here we will learn other ways to produce word vectors
  - ▶ Count-based methods
  - ▶ More efficient neural methods designed just for learning word vectors



# Count-Based Methods

# The Vector Space Model

- Fundamental idea: represent meaning as a vector
- Consider documents as context
- One matrix, two viewpoints
  - Documents represented by their words
  - Words represented by their documents

	...	state	fun	heaven	...
...					
425		0	1	0	
426		3	0	0	
427		0	0	0	
.....					



# Manipulating the VSM

- Weighting the values (beyond frequency)
- Creating low-dimensional dense vectors
- Comparing vectors

# Tf-idf

- Standard weighting scheme for information retrieval

$$idf_w = \log \frac{|D|}{df_w}$$

- Discounts common words

	...	the	country	hell	...
...					
425		43	5	1	
426		24	1	0	
427		37	0	3	
...					
df		500	14	7	

*tf* matrix

	...	the	country	hell	...
...					
425		0	26.0	6.2	
426		0	5.2	0	
427		0	0	18.6	
...					

*tf-idf* matrix

# Dimensionality Reduction

- Term-document matrices are very **sparse**
- Dimensionality reduction: create shorter, denser vectors
- More practical (less features)
- Remove noise (less overfitting)

# Singular Value Decomposition

|D|

|V|

$$\begin{bmatrix} 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 0 & \dots & 0 \\ 0 & 0 & 0 & \dots & 0 \\ \vdots & & & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 0 \end{bmatrix}$$

$A$   
(term-document matrix)

$$A = U \Sigma V^T$$

$U$

(new term matrix)

$m$

$\Sigma$

(singular values)

$m$

$V^T$

(new document matrix)

|D|

|V|

$$\begin{bmatrix} 2.2 & 0.3 & \dots & 8.7 \\ 5.5 & -2.8 & \dots & 0.1 \\ -1.3 & 3.7 & \dots & 3.5 \\ \vdots & & \ddots & \vdots \\ 2.9 & -2.1 & \dots & -1.9 \end{bmatrix}$$

$m$

$$\begin{bmatrix} 9.1 & 0 & 0 & \dots & 0 \\ 0 & 4.4 & 0 & \dots & 0 \\ 0 & 0 & 2.3 & \dots & 0 \\ \vdots & & & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 0.1 \end{bmatrix}$$


$m$

$$\begin{bmatrix} -0.2 & 4.0 & \dots & -1.3 \\ -4.1 & 0.6 & \dots & -0.2 \\ 2.6 & 6.1 & \dots & 1.4 \\ \vdots & & \ddots & \vdots \\ -1.9 & -1.8 & \dots & 0.3 \end{bmatrix}$$

$$m = \text{Rank}(A)$$

# Truncating – Latent Semantic Analysis

- Truncating  $U$ ,  $\Sigma$ , and  $V$  to  $k$  dimensions produces best possible  $k$  rank approximation of original matrix
- So truncated,  $U_k$  (or  $V_k^T$ ) is a new low dimensional representation of the word
- Typical values for  $k$  are 100-5000



# Words as Context

- Lists how often words appear with other words
  - In some predefined context (usually a window)
- The obvious problem with raw frequency: dominated by common words

	...	the	country	hell	...
...					
state		1973	10	1	
fun		54	2	0	
heaven		55	1	3	
.....					

# Pointwise Mutual Information

- For two events  $x$  and  $y$ , PMI computes the discrepancy between:
  - ▶ Their joint distribution
  - ▶ Their individual distributions (assuming independence)

$$PMI(x, y) = \log_2 \frac{p(x, y)}{p(x)p(y)}$$

# Calculating PMI

	...	the	country	hell	...		$\Sigma$
...							
state		1973	10	1			12786
fun		54	2	0			633
heaven		55	1	3			627
...							
$\Sigma$		1047519	3617	780			15871304

$$p(x,y) = \text{count}(x,y) / \Sigma$$

$$p(x) = \Sigma_x / \Sigma$$

$$p(y) = \Sigma_y / \Sigma$$

$x = \text{state}, y = \text{country}$

$$p(x,y) = 10 / 15871304 = 6.3 \times 10^{-7}$$

$$p(x) = 12786 / 15871304 = 8.0 \times 10^{-4}$$

$$p(y) = 3617 / 15871304 = 2.3 \times 10^{-4}$$

$$\begin{aligned} \text{PMI}(x,y) &= \log_2(6.3 \times 10^{-7}) / ((8.0 \times 10^{-4}) (2.3 \times 10^{-4})) \\ &= 1.78 \end{aligned}$$



# PMI Matrix

- PMI does a better job of capturing interesting semantics
  - ▶ E.g. *heaven* and *hell*
- But it is obviously biased towards rare words
- And doesn't handle zeros well

	...	the	country	hell	...
...					
state		1.22	1.78	0.63	
fun		0.37	3.79	-inf	
heaven		0.41	2.80	6.60	
.....					

# PMI Tricks

- Zero all negative values (PPMI)
  - ▶ Avoid  $-\infty$  and unreliable negative values
- Counter bias towards rare events
  - ▶ Smooth probabilities

# SVD

	...	the	country	hell	...
...					
425		0	26.0	6.2	
426		0	5.2	0	
427		0	0	18.6	
...					

*tf-idf* matrix

	...	the	country	hell	...
...					
state		1.22	1.78	0.63	
fun		0.37	3.79	0	
heaven		0.41	2.80	6.60	
.....					

PPMI matrix

- Regardless of whether we use document or word as context, SVD can be applied to create dense vectors

# Similarity

- Word similarity = comparison between word vectors (e.g. cosine similarity)
- Find synonyms, based on proximity in vector space
  - ▶ automatic construction of lexical resources
- Use vectors as features in classifier — more robust to different inputs (*movie* vs *film*)

# Neural Methods

# Word Embeddings

- We've seen word embeddings used in neural networks (feedforward or recurrent)
- But these models are designed for other tasks:
  - ▶ Classification
  - ▶ Language modelling
- Word embeddings is just one part of the model

# Neural Models for Embeddings

- Can we design neural networks whose goal is to learn word embeddings?
- Desiderata:
  - ▶ Unsupervised
  - ▶ Efficient
  - ▶ Useful representation

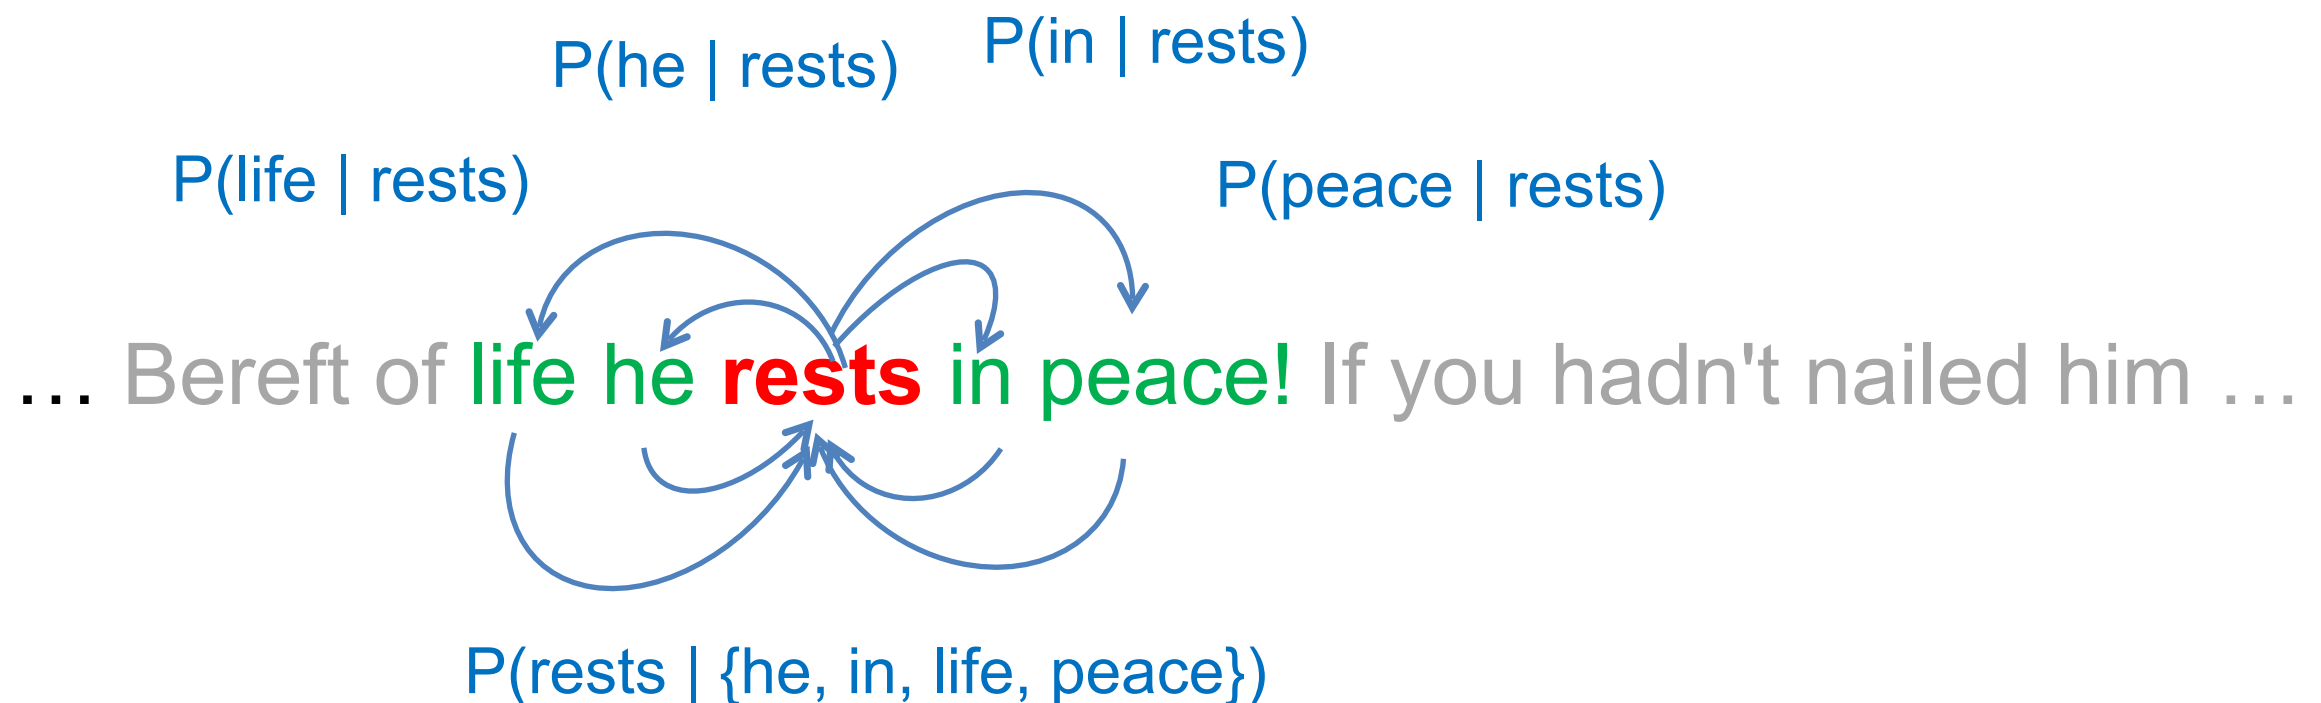
# Word2Vec

- Neural network inspired approaches seek to learn vector representations of words and their contexts
- Key idea
  - ▶ *Word embeddings should be **similar** to embeddings of **neighbouring** words*
  - ▶ *And **dissimilar** to other words that don't occur nearby*
- Using vector dot product for vector 'comparison'
  - ▶  $u \cdot v = \sum_j u_j v_j$



# Word2Vec

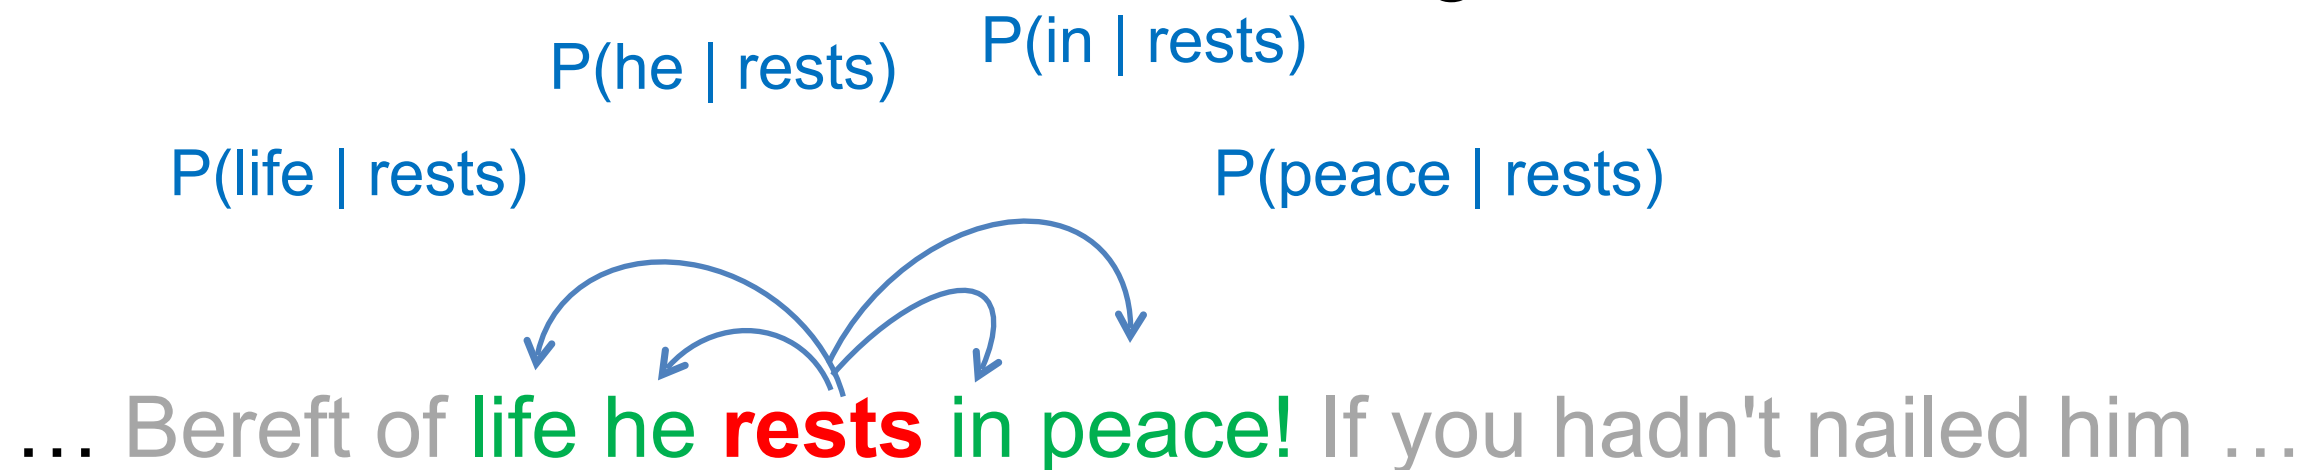
- Framed as learning a classifier
  - ▶ **Skip-gram**: predict words in local context surrounding given word



- ▶ **CBOW**: predict word in centre, given words in the local surrounding context
- Local context means words within  $L$  positions,  $L=2$  above

# Skip-gram Model

- Generates each word in context given centre word



- Total probability defined as

► Where subscript denotes position in running text

$$\prod_{l \in -L, \dots, -1, 1, \dots, L} P(w_{t+l} | w_t)$$

- Using a logistic regression model

$$P(\overset{\text{he}}{w_k} | \overset{\text{rests}}{w_j}) = \frac{\exp(c_{w_k} \cdot v_{w_j})}{\sum_{w' \in V} \exp(c_{w'} \cdot v_{w_j})}$$

# Embedding parameterisation

- Two parameter matrices, with  $d$ -dimensional embedding for all words




Fig 19.17, JM3

- Words are numbered, e.g., by sorting vocabulary and using word location as its index

# Skip-gram model




Fig 19.18, JM3

# Training the skip-gram model

- Train to maximise likelihood of **raw text**
- Too slow in practice, due to normalisation over  $|V|$

$$P(w_k | w_j) = \frac{\exp(c_{w_k} \cdot v_{w_j})}{\sum_{w' \in V} \exp(c_{w'} \cdot v_{w_j})}$$

- Reduce problem to binary classification, distinguish real context words from non-context words aka “**negative samples**”
  - ▶ words drawn randomly from  $V$

# Negative Sampling


... lemon,    a [tablespoon of apricot jam,         a] pinch ...  
                c1                      c2         t         c3                      c4

**positive examples +**

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

## negative examples -

t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if



$$P(+|t, c) = \frac{1}{1 + e^{-t \cdot c}} \leftarrow \text{maximise similarity between target word and real context words}$$

$$P(-|t, c) = 1 - \frac{1}{1 + e^{-t \cdot c}} \leftarrow \text{minimise similarity between target word and non-context words}$$

# Skip-gram Loss

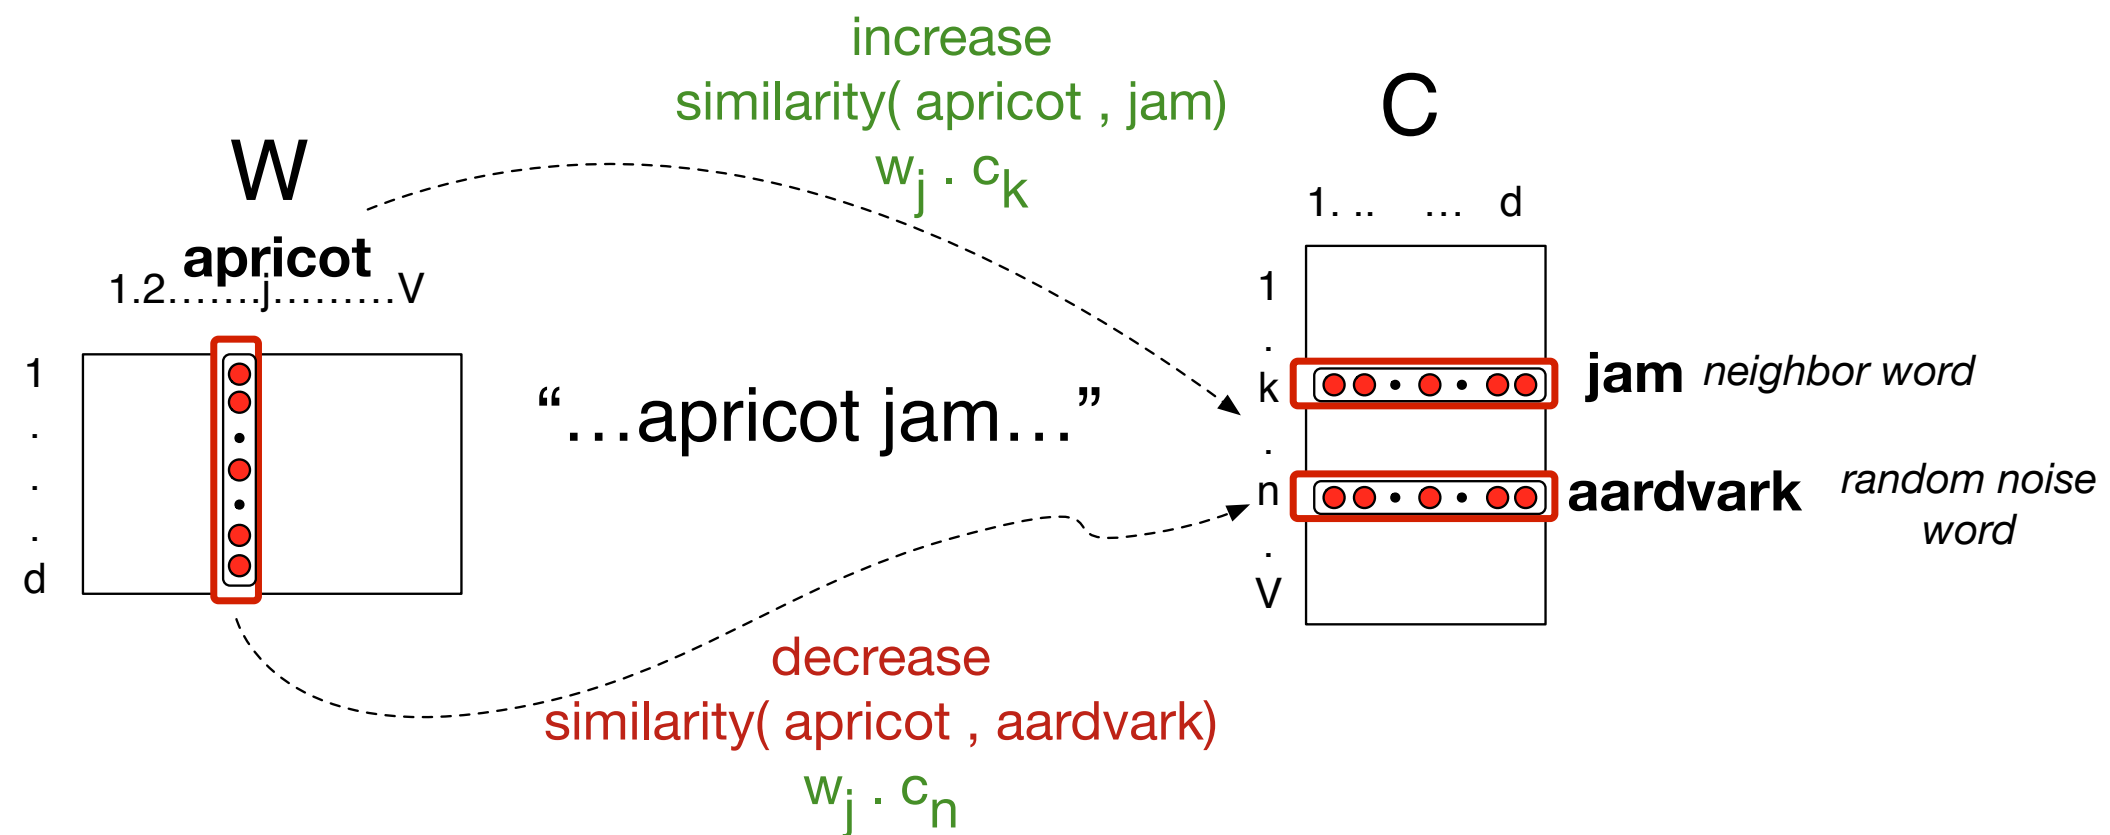
$$L(\theta) = \sum_{(t,c) \in +} \log P(+|t, c) + \sum_{(t,c) \in -} \log P(-|t, c)$$

- In practice, use  $k$  negative examples

$$L(\theta) = \log P(+|t, c) + \sum_{i=1}^k \log P(-|t, n_i)$$

# Training Illustration

- Iterative process (stochastic gradient descent)
  - ▶ each step moves embeddings closer for context words
  - ▶ and moves embeddings apart for noise samples







# Desiderata

- Unsupervised
  - ▶ Raw, unlabelled corpus
- Efficient
  - ▶ Negative sampling (avoid softmax over full vocabulary)
  - ▶ Scales to very very large corpus
- Useful representation:
  - ▶ How do we evaluate word vectors?

# Evaluating Word Vectors

- Lexicon style tasks
  - ▶ *WordSim-353* are pairs of nouns with judged relatedness
  - ▶ *SimLex-999* also covers verbs and adjectives
  - ▶ *TOEFL* asks for closest synonym as multiple choice
  - ▶ ...
- Test compatibility of word pairs using cosine similarity in vector space

# Embeddings Exhibit Meaningful Geometry



- Word analogy task

- ▶ *Man* is to *King* as *Woman* is to ???
- ▶  $v(\text{Man}) - v(\text{King}) = v(\text{Woman}) - v(\text{???})$
- ▶  $v(\text{???}) = v(\text{Woman}) - v(\text{Man}) + v(\text{King})$

# Evaluating Word Vectors

- Best evaluation is in other downstream tasks
  - ▶ Use bag-of-word embeddings as a feature representation in a classifier
  - ▶ First layer of most deep learning models is to embed input text; use pre-trained word vectors as embeddings
- Recently **contextual word vectors** shown to work even better
  - ▶ ELMO & BERT (next lecture!)

# Pointers to Software

- Word2Vec
  - ▶ C implementation of Skip-gram and CBOW  
<https://code.google.com/archive/p/word2vec/>
- GenSim
  - ▶ Python library with many methods include LSI, topic models and Skip-gram/CBOW  
<https://radimrehurek.com/gensim/index.html>
- GLOVE
  - ▶ <http://nlp.stanford.edu/projects/glove/>

# Further Reading

- ▶ JM3, Ch 6