School of Computing and Information Systems
The University of Melbourne
COMP90042
NATURAL LANGUAGE PROCESSING (Semester 1, 2020)

Sample solutions for discussion exercises: Week 6

Discussion

1. Give illustrative examples that show the difference between:

   (a) **Synonyms** and **hypernyms**
      - Two words are synonyms when they share (mostly) the same meaning, for example: *snake* and *serpent* are synonyms.
      - One word is a hypernym of a second word when it is a more general instance ("higher up" in the hierarchy) of the latter, for example, *reptile* is the hypernym of *snake* (in its animal sense).

   (b) **Hyponyms** and **meronyms**
      - One word is a hyponym of a second word when it is a more specific instance ("lower down" in the hierarchy) of the latter, for example, *snake* is one hyponym of *reptile*. (The opposite of hypernymy.)
      - One word is a meronym of a second word when it is a part of the whole defined by the latter, for example, *scales* (the skin structure) is a meronym of *reptile*.

2. Using some Wordnet visualisation tool, for example,

   `http://wordnetweb.princeton.edu/perl/webwn` and the Wu & Palmer definition of **word similarity**, check whether the word *information* is more similar to the word *retrieval* or the word *science* (choose the sense which minimises the distance). Does this mesh with your intuition?

   - The word *information* has five different senses in Wordnet; I've reproduce the fragment of the hierarchy above these senses below:

		entity abstraction... communication message...		entity
entity	entity	statement	entity	abstraction...
abstraction...	abstraction...	pleading	abstraction...	measure
communication	psychological...	charge...	group...	system of meas...
message...	cognition...	accusation...	collection...	information meas...

<div align="center">information</div>

- Here's the corresponding fragment of the three senses above `retrieval`:

```
entity              entity
physical...         abstraction...
process...          psychological...  entity
processing          cognition...      abstraction...
data process... |   process...        psychological...
operation           basic cog...      event
computer op... |    memory...         act...
                    retrieval
```

- To find the Wu & Palmer similarity, we need to find the **lowest common subsumer** — the lowest node in the hierarchy shared by the two senses, and then apply the following formula:

$$\text{sim}(c_1, c_2) = \frac{2 \times \text{depth}(\text{LCS}(c_1, c_2))}{\text{depth}(c_1) + \text{depth}(c_2)}$$

- The question asks us to choose the senses which minimise the distance, so we need to check them all.

- The "message" sense of *information* lies at depth 5; the "data retrieval" sense of *retrieval* is at depth 8; the lowest node in the hierarchy that they share is `entity` (the root node) so the similarity is:

$$
\begin{aligned}
\text{sim}(\text{information}, \text{retrieval}) &= \frac{2 \times 1}{5 + 8} \\
&= \frac{2}{13} \approx 0.154
\end{aligned}
$$

- What about the "message" sense of *information* with the other senses of *retrieval*?

  - The "memory" sense of *retrieval* is at depth 8, and the lowest node shared is `abstraction, abstract entity` (at depth 2); this means that the similarity is $\frac{2 \times 2}{8+5} \approx 0.308$.
  - The "event" sense of *retrieval* is at depth 6, and the lowest node shared is also `abstraction, abstract entity`, so the similarity is $\frac{2 \times 2}{6+5} = 0.364$.

- Let me summarise these results in a table, where I've numbered the senses according to the Wordnet ordering (left-to-right above):

		\multicolumn{5}{c}{*information*}
		1	2	3	4	5
	1	0.154	0.154	0.118	0.154	0.143
retrieval	2	0.308	**0.615**	0.235	0.308	0.286
	3	0.364	0.545	0.267	0.364	0.333

- The maximum similarity (in bold in the table above) is 0.615, for the second sense of *information* — "knowledge acquired through study or experience or instruction" — and the second sense of *retrieval* — "the cognitive operation of accessing information in memory" (because they are both cognitive processes).

- I will leave *science* as an exercise (there are only two senses this time), but the maximum similarity is 0.727 for the "knowledge acquired..." sense of *information*, and the "ability to produce solutions in some problem domain" sense of *science*.
- *science* is clearly the more similar word. This does match with my personal expectations, however, this probably isn't the sense of *science* I had in mind!

3. What is **word sense disambiguation**?

  - Word sense disambiguation is the computational problem of automatically determining which sense (usually, Wordnet synset) of a word is intended for a given token instance with a document.

4. For the following term co-occurrence matrix (suitably interpreted):

	cup	not(cup)
world	55	225
not(world)	315	1405

  (a) Find the Point-wise Mutual Information (PMI) between these two terms in this collection.
    - To evaluate PMI, we need the joint and prior probabilities of the two event (in this case, probably w: document contains `world` and c: document contains `cup`.
    - We estimate these based on their appearance out of the total number of instances in the collection (2000), and then substitute:

$$\begin{aligned}
P(w) &= 280/2000 = 0.14 \\
P(c) &= 370/2000 = 0.185 \\
P(w,c) &= 55/2000 = 0.0275 \\
PMI(w,c) &= \log_2 \frac{P(w,c)}{P(w)P(c)} \\
&= \log_2 \frac{0.0275}{0.14 \times 0.185} \\
&\approx 0.0865
\end{aligned}$$

  (b) What does the value from (a) tell us about **distributional similarity**?
    - This value is slightly positive, which means that the two events occur together (in documents) slightly more commonly than would occur purely by chance. There is some possibility that `world` and `cup` occurring together is somehow meaningful for documents in this collection.

5. In the `09-distributional-semantics` iPython notebook, a term–document matrix is built to learn word vectors.

  (a) What is the Singular Value Decomposition (SVD) method used for here? Why is this helpful?
    - We are using the SVD method to build a representation of our matrix which can be used to identify the most important characteristics of words.
    - By throwing away the less important characteristics, we can have a smaller representation of the words, which will save us (potentially a great deal of) time when evaluating the cosine similarities between word pairs.

6. What is a **word embedding** and how does it relate to **distributional hypothesis**?

- We're going to have a representation of words (based on their contexts) in a **vector space**, such that other words "nearby" in the space are similar
- This is broadly the same as what we expect in distributional similarity ("you shall know a word by the company it keeps.")
- The row corresponding to the word in the relevant (target/context) matrix is known as the "embedding".

(a) What is the difference between a **skip-gram** model and a **CBOW** model?

- The element in the condition of the posterior probability: skip–gram models analyse the probability of the context words **given** the target word; CBOW models analyse the probability of the target word **given** the context words.

(b) How are the above models trained?

- The probabilities here are more complicated than just counting some events in a collection; they are based around taking the dot product of the relevant vectors (or average of vectors, in the case of CBOW), and then **marginalising**. To improve efficiency, we can also use a negative sampling objective.